

This is a specification of the commit-adopt algorithm of Gafni and Losa in the message-adversary model with dynamic participation. The specification is written in PlusCal and TLA+.

The message-adversary model with dynamic participation is like the sleepy model, except that processes never fail; instead, the adversary corrupts their messages. This has the same effect as processes being faulty but is cleaner to model.

Note that, to check this specification with the TLC model-checker, you must first translate the PlusCal algorithm to TLA+ using the TLA toolbox or the TLA+ VSCode extension.

EXTENDS *Naturals*, *FiniteSets*, *Utilities*

CONSTANTS

- $P$  the set of processors
- ,  $V$  the set of possible values
- ,  $Bot$  the special value “bottom”, indicating the absence of something
- ,  $Lambda$  the failure notification “lambda”
- ,  $NoCommit$  an indication that a processors didn’t see a unanimous majority in round 1 of the algorithm

ASSUME  $Distinct(\langle P, V, \{Bot\}, \{Lambda\}, \{NoCommit\} \rangle)$

**--algorithm** *CA*{

**variables**

- $input \in [P \rightarrow V]$ ; the processors’ inputs
- $sent = [p \in P \mapsto Bot]$ ; messages sent in the current round
- message received by  $p$  from  $q$  in the current round;  $Bot$  means no message received:
- $received = [p \in P \mapsto [q \in P \mapsto Bot]]$ ;
- $rnd = 1$ ; the current round (1, 2, or 3); we end at 3 but nothing happens in round 3
- the processors’ outputs (either  $Bot$ ,  $\langle \text{“commit”}, v \rangle$ , or  $\langle \text{“adopt”}, v \rangle$ ) for some  $v$
- $output = [p \in P \mapsto Bot]$ ;

**define** {

- $TypeOkay \triangleq$ 
  - $\wedge input \in [P \rightarrow V]$
  - $\wedge sent \in [P \rightarrow V \cup \{Bot, NoCommit\}]$
  - $\wedge received \in [P \rightarrow [P \rightarrow V \cup \{Bot, NoCommit, Lambda\}]]$
  - $\wedge rnd \in \{1, 2, 3\}$
  - $\wedge output \in [P \rightarrow \{Bot\} \cup \{\langle ca, v \rangle : ca \in \{\text{“commit”}, \text{“adopt”}\}, v \in V\}]$
- the set of processors from which  $p$  received a message (*i.e.* heard of):
- $HeardOf(p) \triangleq \{q \in P : received[p][q] \neq Bot\}$
- the set of minority subsets of  $S$ :
- $Minority(S) \triangleq \{M \in SUBSET S : 2 * Cardinality(M) < Cardinality(S)\}$
- the number of votes for  $v$  that  $p$  received:
- $VoteCount(p, v) \triangleq Cardinality(\{q \in P : received[p][q] = v\})$
- the set of values  $v$  for which  $p$  received a strict majority of votes:
- $VotedByMajority(p) \triangleq \{v \in V : 2 * VoteCount(p, v) > Cardinality(HeardOf(p))\}$
- the set containing the value voted for most often, if any, according to  $p$ ’s received messages:

$MostVotedFor(p) \triangleq \{v \in V : \forall w \in V \setminus \{v\} : VoteCount(p, v) > VoteCount(p, w)\}$

for technical reasons, we need the program counter of a processor in round  $r$ :

$Pc(r) \triangleq \text{CASE } r = 1 \rightarrow \text{"r1"}$

$\square r = 2 \rightarrow \text{"r2"}$

$\square r = 3 \rightarrow \text{"r3"}$

Now we give the two safety properties:

$Agreement \triangleq \forall p, q \in P : output[p] \neq Bot \wedge output[q] \neq Bot \wedge output[p][1] = \text{"commit"}$   
 $\Rightarrow output[p][2] = output[q][2]$

$Validity \triangleq \forall p \in P : \forall v \in V :$

$pc[p] = \text{"Done"} \wedge (\forall q \in P : input[q] = v) \Rightarrow output[p] = \langle \text{"commit"}, v \rangle$

}

**macro** *broadcast*(  $v$  ) {

$sent := [sent \text{ EXCEPT } ![self] = v]$

}

The following macro is used to deliver messages to the processors. It includes message corruptions by the adversary:

**macro** *deliver\_msgs*( *participating*, *corrupted* ) {

**with** (  $ByzMsg \in [P \rightarrow [corrupted \rightarrow V \cup \{Bot, Lambda, NoCommit\}]]$  ) {

we assert the properties of the no-equivocation model:

**when**  $\forall p1, p2 \in P : \forall q \in corrupted :$

$ByzMsg[p1][q] \in V \Rightarrow ByzMsg[p2][q] \in \{ByzMsg[p1][q], Lambda\};$

$received := [p \in P \mapsto [q \in P \mapsto$

IF  $q \in corrupted$

THEN  $ByzMsg[p][q]$   $p$  receives a corrupted message

ELSE IF  $q \in participating$

THEN  $sent[q]$   $p$  receives what  $q$  sent

ELSE  $Bot$ ]]  $p$  receives nothing

}

}

Now we give the specification of the algorithm:

**fair process** (  $proc \in P$  ) {

in round 1, vote for  $input[self]$ :

$r1:$   $broadcast(input[self]);$

$r2:$  **await**  $rnd = 2;$

if there is a majority for a value  $v$ , propose to commit  $v$ :

**if** (  $VotedByMajority(self) \neq \{\}$  )

**with** (  $v \in VotedByMajority(self)$  ) the set is a singleton at this point

$broadcast(v)$

**else**

$broadcast(NoCommit);$

$r3:$  **await**  $rnd = 3;$  in round 3 we just produce an output

**if** (  $VotedByMajority(self) \neq \{\}$  ) if there is a majority for a value  $v$ , commit  $v$ :

**with** (  $v \in VotedByMajority(self)$  ) the set is a singleton at this point

$output[self] := \langle \text{"commit"}, v \rangle$

**else if** (  $MostVotedFor(self) \neq \{\}$  ) otherwise, adopt a most voted value:

```

with (  $v \in \text{MostVotedFor}(\text{self})$  ) there can be multiple values in the set
 $\text{output}[\text{self}] := \langle \text{"adopt"}, v \rangle$ 
else if no value was voted for, adopt an arbitrary value:
with (  $v \in V$  )
 $\text{output}[\text{self}] := \langle \text{"adopt"}, v \rangle$ 
}

```

Below we specify the behavior of the adversary. The no-equivocation model guarantees that if a processor receives  $v$  from  $p$ , then all receive  $v$  or  $\text{Lambda}$ .

```

fair process (  $\text{adversary} \in \{\text{"adversary"}\}$  ) {
adv: while (  $\text{rnd} < 3$  ) {
    await  $\forall p \in P : \text{pc}[p] = \text{Pc}(\text{rnd} + 1)$ ;
    pick a participating set and a set of corrupted processors:
    with (  $\text{Participating} \in \text{SUBSET } P \setminus \{\{\}\}$  )
    with (  $\text{Corrupted} \in \text{Minority}(\text{Participating})$  )
         $\text{deliver\_msgs}(\text{Participating}, \text{Corrupted})$ ;
     $\text{rnd} := \text{rnd} + 1$ ;
}
}
}

```

Canary invariants that should break (this is to make sure that the specification reaches expected states):

To find a state in which some process outputs:

$\text{Canary1} \triangleq \forall p \in P : \text{output}[p] = \text{Bot}$

To find a state in which some process commits while another adopts:

$\text{Canary2} \triangleq \forall p, q \in P :$   
 $\wedge \text{output}[p] \neq \text{Bot}$   
 $\wedge \text{output}[q] \neq \text{Bot}$   
 $\Rightarrow \neg(\text{output}[p][1] = \text{"commit"} \wedge \text{output}[q][1] = \text{"adopt"})$

To find a state in which two processes adopt different values:

$\text{Canary3} \triangleq \forall p, q \in P :$   
 $\wedge \text{output}[p] \neq \text{Bot}$   
 $\wedge \text{output}[q] \neq \text{Bot}$   
 $\Rightarrow \neg(\text{output}[p][1] = \text{"adopt"} \wedge \text{output}[q][1] = \text{"adopt"} \wedge \text{output}[p][2] \neq \text{output}[q][2])$

---

\ \* Modification History  
 \ \* Last modified Sun Jan 01 16:07:58 PST 2023 by nano  
 \ \* Created Thu Dec 29 09:54:34 PST 2022 by nano